

How to find lots of bugs in real code with system-specific static analysis

Dawson Engler
Stanford University

Ben Chelf, Andy Chou, Seth Hallem
Coverity

The core static bug finding intuition

- Systems software has many ad-hoc restrictions:
 - “acquire lock L before accessing shared variable X”
 - “disabled interrupts must be re-enabled”
 - Error = crashed system. How to find?
- Observation: rules can be checked with a compiler
 - scan source for “relevant” acts, check that correct.
 - E.g., to check “disabled interrupts must be re-enabled:” scan for calls to `disable()/enable()`, check matching, not done twice
- Main problem:
 - compiler has machinery to check, but not knowledge
 - implementor has knowledge but not machinery
- System-specific static analysis:
 - give implementors a framework to add easily-written, system-specific compiler extensions

System-specific static analysis

□ Implementation:

Extensions dynamically linked into EDG C compiler

Applied down all paths (“flow sensitive”), across all procedures (“interprocedural”) in input program source at compile time.

Linux
drivers/
raid5.c

```
save(flags);  
cli();  
if(!(buf = kmalloc()))  
    return 0;  
restore(flags);  
return buf;
```

EDG compiler

int checker

“did not re-
enable ints!”

Scalable: handles millions of lines of code

Precise: says exactly what error was


Immediate: finds bugs without having to execute path

Effective: 1000s of bugs in Linux, OpenBSD, Commercial

A bit more detail

```
{ #include "linux-includes.h" }
sm chk_interrupts {
  decl { unsigned } flags;
  // named patterns
  pat enable = { sti(); }
             | { restore_flags(flags); };
  pat disable = { cli(); };


  // states
  is_enabled: disable ==> is_disabled
    | enable ==> { err("double enable"); }
    ;
  is_disabled: enable ==> is_enabled
    | disable ==> { err("double disable"); }
    | $end_of_path$ ==>
      { err("exiting w/intr disabled!"); }
    ; }
```



No X after Y: do not use freed memory

```
sm free_checker {  
  state decl any_pointer v;  
  decl any_pointer x;  
  
  start: { kfree(v); } ==> v.freed  
  ;  
  v.freed:  
    { v != x } || { v == x }  
      ==> { /* do nothing */ }  
  | { v } ==> { err("Use after free!"); }  
  ;  
}
```

```
/* 2.4.1: fs/proc/generic.c */  
ent->data = kmalloc(...)  
if(!ent->data) {  
    kfree(ent);  
    goto out;  
...  
out: return ent;
```



"In context Y, don't do X": blocking

- Linux: if interrupts are disabled, or spin lock held, do not call an operation that could block:


Compute transitive closure of all potentially blocking fn's

Hit disable/lock: warn of any calls

123 errors, 8 false pos

```
drivers/net/pcmcia/wavelan_cs.c */
n_lock_irqsave (&lp->lock, flags); /* 1889 */
atch(cmd)

case SIOCGIWPRIV: /* 2304 */
if(copy_to_user(wrq->u.data.pointer, ...))
```

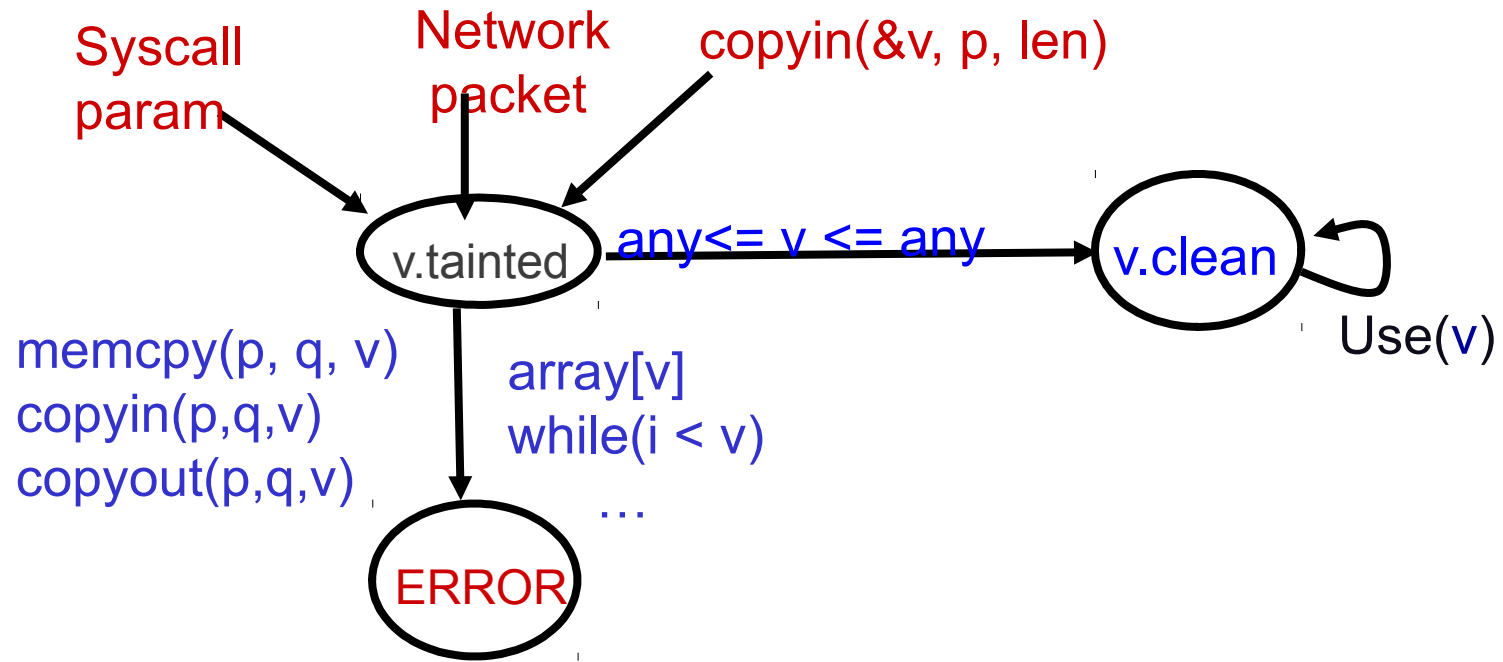


Heavy clustering:

net/atm: 152 checks, 22 bugs (exp 1.9) $P = 3.1 \times 10^{-15}$

"X before Y": sanitize integers before use

- Security: OS must check user integers before use
- Checker: Warn when unchecked integers from **untrusted sources** reach **trusting sinks**



Global; simple to retarget (text file with 2 srcs&12 sinks)
Linux: 125 errors, 24 false; BSD: 12 errors, 4 false

Some big, gaping security holes.

Remote exploit, no checks

```
/* 2.4.9/drivers/isdn/act2000/capi.c:actcapi_dispatch */
isdn_ctrl cmd;
...
while ((skb = skb_dequeue(&card->rcvq))) {
    msg = skb->data;
    ...
    memcpy(cmd.parm.setup.phone, msg->msg.connect_ind.addr.num,
           msg->msg.connect_ind.addr.len - 1);
}
```

Missed lower-bound check:

```
/* 2.4.5/drivers/char/drm/i810_dma.c */
if(copy_from_user(&d, arg, sizeof(arg)))
    return -EFAULT;
if(d.idx > dma->buf_count)
    return -EINVAL;
buf = dma->buflist[d.idx];
Copy_from_user(buf_priv->virtual, d.address, d.used);
```


Enforcing subtle rules

Unexpected overflow

```
copy_from_user(&wrthdr, addr, sizeof wrthdr);  
if ( wrthdr.size + wrthdr.offset > FST_MEMSIZE )  
    return -ENXIO;  
copy_from_user(card->mem+wrthdr.offset,data,wrthdr.size)
```

```
/* 2.4.9-ac7/fs/intermezzo/psdev.c */  
error = copy_from_user(&input, arg, sizeof(input));  
input.path = kmalloc(input.path_len + 1, GFP_KERNEL);  
if ( !input.path )  
    return -ENOMEM;  
error =copy_from_user(input.path,user_path, input.path_len);
```

Weird security implications

```
get_user(len, oldlenp); /* 2.4.1/kernel/sysctl.c */  
if (len > table->maxlen)  
    len = table->maxlen;  
copy_to_user(oldval, table->data, len);
```

Results for BSD 2.8 & 4 months of Linux

All bugs released to implementors; most serious fixed

Violation	Linux		BSD	
	Bug Fixed		Bug Fixed	
Gain control of system	18	15	3	3
Corrupt memory	43	17	2	2
Read arbitrary memory	19	14	7	7
Denial of service	17	5	0	0
Minor	28	1	0	0
Total	125	52	12	12

Local bugs	109	12
Global bugs	16	0
Bugs from inferred ints	12	0
False positives	24	4
Number of checks	~3500	594

Talk Overview

▮ System-specific static analysis:

Correctness rules map clearly to concrete source actions

Check by making compilers aggressively system-specific

Easy: digest sentence fragment, write checker.

One person writes checker, imposed on all code.

Result: precise, immediate error diagnosis. Found errors in every system looked at

▮ Next: Belief analysis

Using programmer beliefs to infer state of system, relevant rules

Key: Find bugs without knowing truth.

Goal: find as many serious bugs as possible

Problem: what are the rules?!?!?

100-1000s of rules in 100-1000s of subsystems.

To check, must answer: Must `a()` follow `b()`? Can `foo()` fail? Does `bar(p)` free `p`? Does lock `l` protect `x`?

Manually finding rules is hard. So don't. Instead infer what code believes, cross check for contradiction

Intuition: how to find errors without knowing truth?

Contradiction. To find lies: cross-examine. Any contradiction is an error.

Deviance. To infer correct behavior: if 1 person does X, might be right or a coincidence. If 1000s do X and 1 does Y, probably an error.

Crucial: we know contradiction is an error without knowing the correct belief!

Cross-checking program belief systems

MUST beliefs:

Inferred from acts that imply beliefs code **must** have.

```
x = *p / z; // MUST belief: p not null
           // MUST: z != 0
unlock(l);  // MUST: l acquired
x++;        // MUST: x not protected by l
```

Check using internal consistency: infer beliefs at different locations, then cross-check for contradiction

MAY beliefs: could be coincidental

Inferred from acts that imply beliefs code **may** have

```
A(); A(); A();
..... .. // MAY: A() and B()
B(); B(); B(); // must be paired
```

B(); // MUST: B() need not
// be preceded by A()

Trivial consistency: NULL pointers

- `*p` implies MUST belief:

`p` is not null
- A check (`p == NULL`) implies two MUST beliefs:
 POST: `p` is null on true path, not null on false path
 PRE: `p` was unknown before check
- Cross-check these for three different error types.
- Check-then-use (79 errors, 26 false pos)

```
/* 2.4.1: drivers/isdn/svmb1/capidrv.c */  
if (!card)  
    printk(KERN_ERR, "capidrv-%d: ...", card->contrnr...)
```

Null pointer fun

□ Use-then-check: 102 bugs, 4 false

```
/* 2.4.7: drivers/char/mxser.c */
struct mxser_struct *info = tty->driver_data;
unsigned flags;
if(!tty || !info->xmit_buf)
    return 0;
```

□ Contradiction/redundant checks (24 bugs, 10 false)

```
/* 2.4.7/drivers/video/tdfxfb.c */
fb_info.regbase_virt = ioremap_nocache(...);
if(!fb_info.regbase_virt)
    return -ENXIO;
fb_info.bufbase_virt = ioremap_nocache(...);
if(!fb_info.regbase_virt) {
    iounmap(fb_info.regbase_virt);
```

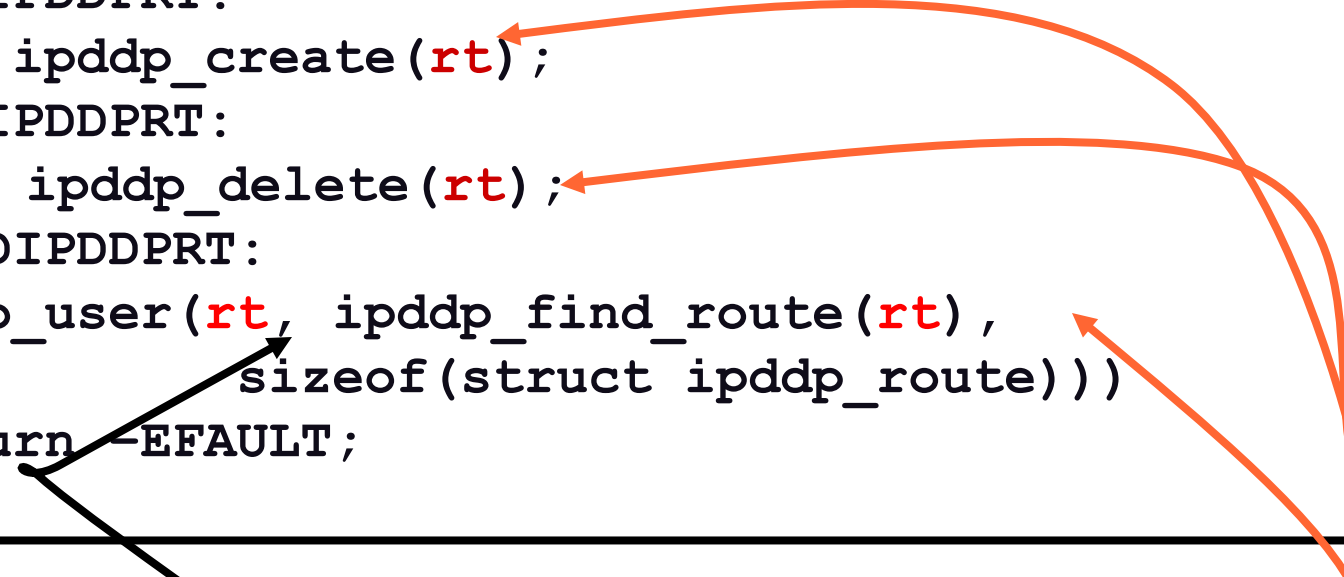
Internal Consistency: finding security holes

- Applications are bad:
 - Rule: “do not dereference user pointer $\langle p \rangle$ ”
 - One violation = security hole
 - Detect with static analysis if we knew which were “bad”
 - Big Problem: which are the user pointers???
- Sol’n: for all pointers, cross-check two OS beliefs
 - “ $\ast p$ ” implies safe kernel pointer
 - “ $\text{copyin}(p)/\text{copyout}(p)$ ” implies dangerous user pointer
 - Error: pointer p has both beliefs.
 - Implemented as a two pass global checker
- Result: 24 security bugs in Linux, 18 in OpenBSD
(about 1 bug to 1 false positive)

An example

Still alive in linux 2.4.4:

```
/* drivers/net/appletalk/ipddp.c:ipddp_ioctl */
case SIOCADDIPDDPRT:
    return ipddp_create(rt);
case SIOCDELIPDDPRT:
    return ipddp_delete(rt);
case SIOFCINDIPDDPRT:
    if(copy_to_user(rt, ipddp_find_route(rt),
        sizeof(struct ipddp_route)))
        return -EFAULT;
```



Tainting marks "rt" as a tainted pointer, checker warns that rt is passed to a routine that dereferences it 2 other examples in same routine...

Cross checking beliefs related abstractly

- Common: multiple implementations of same interface.
Beliefs of one implementation can be checked against those of the others!
- User pointer (3 errors):
If one implementation **taints** its argument, all others must

```
foo_write(void *p, void *arg,...) { bar_write(void *p, void *arg,...) {  
  copy_from_user(p, arg, 4);      *p = *(int *)arg;  
  disable();                     ... do something ...  
  ... do something ...           disable();  
  enable();                      return 0;  
  return 0;                      }  
}
```

How to tell? Routines assigned to same function pointer

```
write_fp = foo_write;  
write_fp = bar_write;
```

MAY beliefs

- ▮ Separate fact from coincidence? General approach:
 - Assume MAY beliefs are MUST beliefs.
 - Check them
 - Count number of times belief passed check (S =success)
 - Count number of times belief failed check (F =fail)
 - Expect: valid beliefs = high ratio of S to F .
 - Use S and F to compute confidence that belief is valid.
 - Rank errors based on this confidence.
 - Go down list, inspecting until false positives are too high.
- ▮ How to weigh evidence?

How to weigh MAY beliefs

- ❑ Wrong way: percentage. (Ignores population size)
Success=1, Failure=0, Percentage = $1/1 * 100 = 100\%$
Success=999, Failure=10, Percentage = $999/1000 = 99.9\%$
- ❑ A better way: "hypothesis testing."
Treat each check as independent binary coin toss
Pick probability p_0 that coin "coincidentally" comes up S.
For a given belief, compute how "unlikely" that it
coincidentally got S successes out of N ($N=S+F$) attempts
$$Z = (\text{observed} - \text{expected}) / \text{stderr}$$
$$= (S - N * p_0) / \sqrt{N * p_0 * (1 - p_0)}$$
- ❑ HUGE mistake: pick T, where $Z > T$ implies MUST
Becomes very sensitive to T.

Statistical: Deriving deallocation routines

- Use-after free errors are horrible.

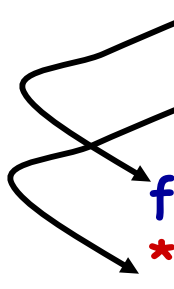
Problem: lots of undocumented sub-system free functions

Soln: derive behaviorally: pointer "p" not used after call "foo(p)" implies MAY belief that "foo" is a free function

- Conceptually: Assume all functions free all arguments
(in reality: filter functions that have suggestive names)

Emit a "check" message at every call site.

Emit an "error" message at every use



foo(p);	foo(p);	foo(p);	bar(p);	bar(p);	bar(p);
*p = x;	*p = x;	*p = x;	p = 0;	p = 0;	*p = x;

Rank errors using z test statistic: $z(\text{checks}, \text{errors})$

E.g., $\text{foo.z}(3, 3) < \text{bar.z}(3, 1)$ so rank bar's error first

Results: 23 free errors, 11 false positives

Ranked free errors

```
kfree[0]: 2623 checks, 60 errors, z= 48.87
  2.4.1/drivers/sound/sound_core.c:sound_insert_unit:
    ERROR:171:178: Use-after-free of 's'! set by 'kfree'
  ...
kfree_skb[0]: 1070 checks, 13 errors, z = 31.92
  2.4.1/drivers/net/wan/comx-proto-fr.c:fr_xmit:
    ERROR:508:510: Use-after-free of 'skb'! set by 'kfree_skb'
  ...
[FALSE] page_cache_release[0] ex=117, counter=3, z = 10.3
dev_kfree_skb[0]: 109 checks, 4 errors, z=9.67
  2.4.1/drivers/atm/iphase.c:rx_dle_intr:
    ERROR:1321:1323: Use-after-free of 'skb'! set by 'dev_kfree_skb_any'
  ...
cmd_free[1]: 18 checks, 1 error, z=3.77
  2.4.1/drivers/block/cciss.c:667:cciss_ioctl:
    ERROR:663:667: Use-after-free of 'c'! set by 'cmd_free[1]'
drm_free_buffer[1] 15 checks, 1 error, z = 3.35
  2.4.1/drivers/char/drm/gamma_dma.c:gamma_dma_send_buffers:
    ERROR:Use-after-free of 'last_buf'!
[FALSE] cmd_free[0] 18 checks, 2 errors, z = 3.2
```

Recall: deterministic free checker

```
sm free_checker {  
  state decl any_pointer v;  
  decl any_pointer x;  
  
  start: { kfree(v); } ==> v.freed  
  ;  
  v.freed:  
    { v != x } || { v == x }  
      ==> { /* do nothing */ }  
  | { v } ==> { err("Use after free!"); }  
  ;  
}
```

A statistical free checker

```
sm free_checker local {
  state decl any_pointer v;
  decl any_fn_call call;
  decl any_pointer x;

  start: { call(v) } ==> v.freed,
    {
      v.data = call.name();
      printf("checking [POP=%s]", v.data);
    }
  ;

  v.freed:
    { v != x } || { v == x } ==> { /* do nothing */ }
  | { v } ==> { err("Use after free! [FAIL=%s]", v.data); }
  ;
}
```


A bad free error

```
/* drivers/block/cciss.c:cciss_ioctl */
if (ioccommand.Direction == XFER_WRITE) {
    if (copy_to_user(...)) {
        cmd_free(NULL, c);
        if (buff != NULL) kfree(buff);
        return( -EFAULT);
    }
}
if (ioccommand.Direction == XFER_READ) {
    if (copy_to_user(...)) {
        cmd_free(NULL, c);
        kfree(buff);
    }
}
cmd_free(NULL, c);
if (buff != NULL) kfree(buff);
```

Deriving "A() must be followed by B()"


- "a(); ... b();" implies *MAY* belief that a() follows b()
Programmer may believe a-b paired, or might be a coincidence.

- Algorithm:

Assume every a-b is a valid pair (reality: prefilter functions that seem to be plausibly paired)

Emit "success" for each path that has a() then b()

Emit "error" for each path that has a() and no b()



Rank errors for each pair using the test statistic
 $z(\text{foo.success}, \text{foo.error}) = z(2, 1)$

- Results: 23 errors, 11 false positives.

Checking derived lock functions

Evilest: `/* 2.4.1: drivers/sound/trident.c:`
`trident_release:`

```
lock_kernel();  
card = state->card;  
dmabuf = &state->dmabuf;  
VALIDATE_STATE(state);
```

And the award for best effort:

```
/* 2.4.0:drivers/sound/cmpci.c:cm_midi_release: */  
lock_kernel();  
if (file->f_mode & FMODE_WRITE) {  
    add_wait_queue(&s->midi.owait, &wait);  
    ...  
    if (file->f_flags & O_NONBLOCK) {  
        remove_wait_queue(&s->midi.owait, &wait);  
        set_current_state(TASK_RUNNING);  
        return -EBUSY;  
    }  
    ... unlock_kernel();
```

Statistical: deriving routines that can fail

Traditional:

Use global analysis to track which routines return NULL

Problem: false positives when pre-conditions hold, difficult to tell statically ("return p->next"?)

Instead: see how often programmer checks.

Rank errors based on number of checks to non-checks.

Algorithm: Assume *all* functions can return NULL

If pointer checked before use, emit "check" message

If pointer used before check, emit "error"

<code>P = foo(...);</code>	<code>p = bar(...);</code>	<code>p = bar(...);</code>	<code>p = bar(...);</code>	<code>p = bar(...);</code>
<code>*p = x;</code>	<code>If(!p) return;</code>	<code>If(!p) return;</code>	<code>If(!p) return;</code>	<code>*p = x;</code>
	<code>*p = x;</code>	<code>*p = x;</code>	<code>*p = x;</code>	

Sort errors based on ratio of checks to errors

Result: 152 bugs, 16 false.

The worst bug

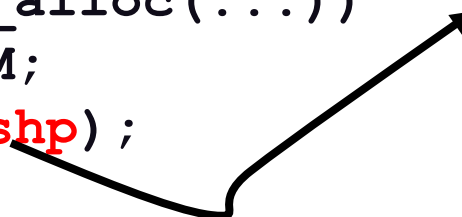
- Starts with weird way of checking failure:

```
/* 2.3.99: ipc/shm.c:1745:map_zero_setup */  
if (IS_ERR(shp = seg_alloc(...)))  
    return PTR_ERR(shp);
```

```
static inline long IS_ERR(const void *ptr)  
{ return (unsigned long)ptr > (unsigned long)-1000L; }
```

- So why are we looking for "seg_alloc"?

```
/* ipc/shm.c:750:newseg: */  
if (!(shp = seg_alloc(...)))  
    return -ENOMEM;  
id = shm_addid(shp);
```



```
int ipc_addid(...* new...) {  
    ...  
    new->cuid = new->uid = ...;  
    new->gid = new->cgid = ...  
    ids->entries[id].p = new;
```

Summary

- ▮ Effective static analysis of real code
 - Write small extension, apply to code, find 100s-1000s of bugs in real systems
 - Result: Static, precise, immediate error diagnosis
 - One person writes, imposes on all code.
- ▮ Belief analysis: broader checking
 - Using programmer beliefs to infer state of system, relevant rules
 - Key feature: find errors without knowing truth
- ▮ Found lots of serious bugs everywhere.

Assertion: Soundness is often a distraction

- Soundness: Find all bugs of type X.
Not a bad thing. More bugs good.
BUT: can only do if you check weak properties.
- What soundness really wants to be when it grows up:
Total correctness: Find all bugs.
Most direct approximation: find as many bugs as possible.
- Opportunity cost:
Diminishing returns: Initial analysis finds most bugs
Spend time on what gets the next biggest set of bugs
Easy experiment: bug counts for sound vs unsound tools.
- Soundness violates end-to-end argument:
“It generally does not make much sense to reduce the residual error rate of one system component (property) much below that of the others.”


Static vs dynamic bug finding

- ▮ Static: precondition = compile (some) code.
All paths + don't need to run + easy diagnosis.
Low incremental cost per line of code
Can get results in an afternoon.
10-100x more bugs.
- ▮ Dynamic: precondition = compile all code + run
What does code do? How to build? How to run?
Runs code, so can check implications.
Good: Static detects ways to cause error, dynamic can check for the error itself.
- ▮ Result:
Static better at checking properties visible in source,
dynamic better at properties implied by source.


Open Q: how to get the bugs that matter?

- Myth: all bugs matter and all will be fixed
 FALSE
 Find 10 bugs, all get fixed. Find 10,000...
- Reality
 All sites have many open bugs (observed by us & PREfix)
 Myth lives because state-of-art is so bad at bug finding
 What users really want: The 5-10 that "really matter"
- General belief: bugs follow 90/10 distribution
 Out of 1000, 100 (10? or 1?) account for most pain.
 Fixing 900+ waste of resources & may make things worse
- How to find worst? No one has a good answer to this.
 Possibilities: promote bugs on executed paths or in code people care about, ...


Open Q: Do static tools really help?



The optimistic hope



The null hypothesis



An Ugly Possibility

Danger: Opportunity cost.

Danger: Deterministic canary bugs to non-deterministic.

Laws of static bug finding

- Vacuous tautologies that imply trouble
 - Can't find code, can't check.
 - Can't compile code, can't check.
- A nice, balancing empirical tautology
 - If can find code
 - AND checked system is big
 - AND can compile (enough) of it
 - THEN: will *always* find serious errors.
- A nice special case:
 - Check rule never checked? Always find bugs. Otherwise immediate kneejerk: what wrong with checker???